

Static Analysis Using Facebook Infer to Find Atomicity Violations

Dominik Harmim

Supervisor: prof. Ing. Tomáš Vojnar, Ph.D.

xharmi00@stud.fit.vutbr.cz

Brno University of Technology, Faculty of Information Technology



24th June 2019

- **Atomicity**: property that a series of operations should be indivisible.
 - Often required in **concurrent programs**.
 - Violation may cause **significant damage**.


```
void replace(int *array, int a, int b) {  
    int i = index_of(array, a);  
    if (i >= 0) set(array, i, b);  
}
```

index_of and **set** should be
executed atomically

The index may be **outdated** because of
a **concurrent modification** of an array.

- **Shortcomings** of current analysers for finding **atomicity violations**:
 - a high rate of **false positives**,
 - **scalability**,
 - ...

- An open-source **static analysis framework** for **interprocedural analyses**.
 - Based on **abstract interpretation**.
- **Highly scalable**.
 - Follows principles of **compositionality**.
 - Computes function **summaries** bottom-up on call trees.
- Supports Java, C, C++ and Objective-C.



- Atomicity violations for sequences of functions.
- Sequences executed atomically once should be executed always atomically.
- Targets C/C++ programs that use PThread locks.

1 Detection of **atomic sequences**.

- Working with **sequences of calls**.
- Only **first occurrences** remembered to assure a termination.

```
void f(int *array) {  
    pthread_mutex_lock(&lock);  
    int i = index_of(array, 42);  
    if (i >= 0) set(array, i, 3);  
    pthread_mutex_unlock(&lock);  
}
```

summary_f: {(index_of, set)}

2 Detection of **violations**.

- Split to sets of **pairs of subsequent calls** without and with a lock.

```
void g(int *array) {  
    int i = index_of(array, 66);  
    if (i >= 0) set(array, i, 5);  
}
```

ATOMICITY VIOLATION!

- The **correctness** was verified on **hand-crafted** programs.
- Evaluated on **real-life low-level concurrent** programs from Debian GNU Linux.
 - Analysed thousands of lines of code.
 - So far quite some false alarms and limited scalability.

Program	Lines of Code	Atomicity Violations
glfw 2.7.9	10,230	13
alsa-utils 1.1.0	7,735	1
c-icap 0.4.2	24,923	174
npth 1.2	1,593	26
rt-tests 0.96	1,795	0
sslsplit 0.4.11	22,457	344

The **static analyser** for finding **atomicity violations**:

- Proposed and implemented as a module of Facebook Infer.
- Successfully tested and experimentally evaluated.
- The preliminary results were presented at Excel@FIT'19.

Future goals:

- Increase the **accuracy**.
 - By, e.g., considering function parameters.
- Enhance the **scalability**.
 - To be able to analyse more extensive programs.
- Create a Pull Request to the `master` branch of Facebook Infer.

Diskutujte obtížnost rozšíření analyzátoru Atomer o podporu formálních parametrů funkcí a návratových hodnot.

- Využití: rozlišení kontextu volaných funkcí.
- Vhodná charakterizace hodnot staticky a abstraktně.
 - Možnost použít syntaktické „access paths“ (lokace na haldě).
 - K bližšímu rozlišení nutno použít ukazatelové analýzy.