



BRNO UNIVERSITY OF TECHNOLOGY

VYSOKÉ UČENÍ TECHNICKÉ V BRNĚ

FACULTY OF INFORMATION TECHNOLOGY

FAKULTA INFORMAČNÍCH TECHNOLOGIÍ

DEPARTMENT OF INTELLIGENT SYSTEMS

ÚSTAV INTELIGENTNÍCH SYSTÉMŮ

STATIC ANALYSIS USING FACEBOOK INFER TO FIND ATOMICITY VIOLATIONS

STATICKÁ ANALÝZA V NÁSTROJI FACEBOOK INFER ZAMĚŘENÁ NA DETEKCI PORUŠENÍ ATOMIČNOSTI

BACHELOR'S THESIS

BAKALÁŘSKÁ PRÁCE

AUTHOR

AUTOR PRÁCE

DOMINIK HARMIM

SUPERVISOR

VEDOUCÍ PRÁCE

prof. Ing. TOMÁŠ VOJNAR, Ph.D.

BRNO 2019

Zadání bakalářské práce



21689

Student: **Harmim Dominik**
Program: Informační technologie
Název: **Statická analýza v nástroji Facebook Infer zaměřená na detekci porušení atomičnosti**
Static Analysis Using Facebook Infer to Find Atomicity Violations
Kategorie: Analýza a testování softwaru

Zadání:

1. Prostudujte principy statické analýzy založené na abstraktní interpretaci. Zvláštní pozornost věnujte přístupům zaměřeným na odhalování problémů v synchronizaci paralelních procesů.
2. Seznamte se s nástrojem Facebook Infer, jeho podporou pro abstraktní interpretaci a s existujícími analyzátory vytvořenými v prostředí Facebook Infer.
3. V prostředí Facebook Infer navrhnete a naimplementujete analyzátor zaměřený na odhalování chyb typu porušení atomičnosti.
4. Experimentálně ověřte funkčnost vytvořeného analyzátoru na vhodně zvolených netriviálních programech.
5. Shrňte dosažené výsledky a diskutujte možnosti jejich dalšího rozvoje v budoucnu.

Literatura:

- Nielson, F., Nielson, H.R., Hankin, C.: Principles of Program Analysis, Springer-Verlag, 2005.
- Blackshear, S., O'Hearn, P.: Open-Sourcing RacerD: Fast Static Race Detection at Scale, 2017. Dostupné on-line: <https://code.fb.com/android/open-sourcing-racerd-fast-static-race-detection-at-scale/>.
- Atkey, R., Sannella, D.: ThreadSafe: Static Analysis for Java Concurrency, Electronic Communications of the EASST, 72, 2015.
- Bielik, P., Raychev, V., Vechev, M.T.: Scalable Race Detection for Android Applications, In: Proc. of OOPSLA'15, ACM, 2015.
- Dias, R.J., Ferreira, C., Fiedor, J., Lourenço, J.M., Smrčka, A., Sousa, D.G., Vojnar, T.: Verifying Concurrent Programs Using Contracts, In: Proc. of ICST'17, IEEE, 2017.

Pro udělení zápočtu za první semestr je požadováno:

- Body 1, 2 a alespoň začátek návrhu z bodu 3.

Podrobné závazné pokyny pro vypracování práce viz <http://www.fit.vutbr.cz/info/szz/>

Vedoucí práce: **Vojnar Tomáš, prof. Ing., Ph.D.**

Vedoucí ústavu: Hanáček Petr, doc. Dr. Ing.

Datum zadání: 1. listopadu 2018

Datum odevzdání: 15. května 2019

Datum schválení: 1. listopadu 2018

Abstract

The goal of this thesis is to propose a *static analyser*, which detects *atomicity violations*. The proposed analyser — *Atomer* — is implemented as an extension for *Facebook Infer*, which is an open-source and extendable static analysis framework that promotes efficient *modular* and *incremental* analysis. The analyser works on the level of *sequences of function calls*. The proposed solution is based on the assumption that sequences executed *once atomically* should probably be executed *always atomically*. The implemented analyser has been successfully verified and evaluated on both smaller programs created for this purpose as well as publicly available benchmarks derived from *real-life low-level* programs.

Abstrakt

Cílem této práce je navrhnout *statický analyzátor*, který bude sloužit pro detekci *porušení atomicity*. Navržený analyzátor — *Atomer* — je implementován jako rozšíření pro *Facebook Infer*, což je volně šířený a snadno rozšiřitelný nástroj, který umožňuje efektivní *modulární* a *inkrementální* analýzu. Analyzátor pracuje na úrovni *sekvencí volání funkcí*. Navržené řešení je založeno na předpokladu, že sekvence, které jsou *jednou zavolány atomicky*, by měly být pravděpodobně volány *atomicky vždy*. Implementovaný analyzátor byl úspěšně ověřen a vyhodnocen jak na malých programech, vytvořených pro tento účel, tak na veřejně dostupných testovacích programech, které vznikly ze *skutečných nízkourovňových* programů.

Keywords

static analysis, programs analysis, abstract interpretation, Facebook Infer, atomicity violation, concurrent programs, contracts for concurrency, atomic sequences, atomicity, incremental analysis, modular analysis, compositional analysis, interprocedural analysis

Klíčová slova

statická analýza, analýza programů, abstraktní interpretace, Facebook Infer, porušení atomicity, paralelní programy, kontrakty pro souběžnost, atomické sekvence, atomicita, inkrementální analýza, modulární analýza, kompoziční analýza, interprocedurální analýza

Reference

HARMIM, Dominik. *Static Analysis Using Facebook Infer to Find Atomicity Violations*. Brno, 2019. Bachelor's thesis. Brno University of Technology, Faculty of Information Technology. Supervisor prof. Ing. Tomáš Vojnar, Ph.D.

Rozšířený abstrakt

Softwarové chyby jsou nedílnou součástí počítačových programů od samotného vzniku programování. Naneštěstí jsou často ukryty na nečekaných místech a mohou vést k nečekanému chování, které může způsobit značné škody. Dnes existuje mnoho způsobů, jak odhalovat tyto chyby již při vývoji. Často se používají *dynamické analyzátoři* nebo nástroje pro automatizované testování. V mnoha případech jsou tyto metody dostačující, nicméně stále mohou zanechat spoustu chyb neodhalených, protože jsou schopny analyzovat pouze určité toky běhu programu na základě vstupních dat. Alternativním řešením je *statická analýza*, která má však také svoje nedostatky. Největším problémem je *škálovatelnost* při analýze rozsáhlých projektů a značně vysoká míra nesprávně hlášených chyb (často se používá anglický výraz „*false alarm*“).

Firma Facebook nedávno představila *Facebook Infer* – nástroj pro tvorbu *vysoce škálovatelných, kompozičních, inkrementálních a interprocedurálních* statických analyzátorů. Facebook Infer značně rozšířil své možnosti, ale je stále v aktivním vývoji mnoha týmy po celém světě. Je používán každý den nejen v samotné firmě Facebook, ale také v jiných firmách jako např. Spotify, Uber, Mozilla nebo Amazon. Momentálně Facebook Infer nabízí několik analyzátorů implementovaných jako moduly v celém aplikačním rámci. Tyto analyzátoři detekují celou řadu typů softwarových chyb, jako např. chyby typu „buffer overflow“ (přetečení vyrovnávací paměti), „thread-safety“ (vláknová bezpečnost), „null-dereferencing“ (dereference prázdného ukazatele) nebo „memory leak“ (únik paměti). Ale především je Facebook Infer *aplikační rámec* pro rychlou a jednoduchou tvorbu nových analyzátorů. V aktuální verzi nástroje Facebook Infer naneštěstí stále chybí lepší podpora pro detekci chyb v *paralelních programech*. Přestože Facebook Infer nabízí docela pokročilé analyzátoři na detekci chyb typu „*data race*“, jsou tyto analyzátoři limitovány pouze na programy napsané v jazyce Java a nejsou navrženy na programy napsané v jazyce C, což by vyžadovalo *nízkoúrovňovější* manipulaci se zámkami.

V paralelních programech se často vyskytují *požadavky na atomicitu* pro provádění určitých sekvencí instrukcí. Porušení těchto požadavků může způsobit různé problémy, jako např. neočekávané chování, výjimky, „segmentation fault“ (chyba porušení paměti) nebo jiné selhání. *Porušení atomicity* obvykle není ověřováno překladačem, na rozdíl od syntaktických nebo některých druhů sémantických pravidel. Požadavky na atomicitu navíc většinou ani vůbec nejsou dokumentovány. Takže v konečném důsledku musí samotní programátoři dbát na dodržování těchto požadavků, a to obvykle bez jakýchkoliv podpůrných nástrojů. Obecně je náročné vyvarovat se těchto chyb v *atomicky závislých programech*, obzvláště ve velkých projektech, a ještě těžší a časově náročnější je hledání a opravování těchto chyb.

V této práci je navržen statický analyzátor — *Atomer* — pro hledání chyb typu porušení atomicity, který je implementován jako rozšíření pro nástroj Facebook Infer. Návrh se konkrétně zaměřuje na *atomické provádění sekvencí volání funkcí*, což je často vyžadováno, např. při použití určitých knihovnických volání. Navržený princip je založen na předpokladu, že sekvence provedené *jednou atomicky*, by pravděpodobně měly být provedeny *atomicky vždy*. Návrh je dále založen na konceptu *kontraktů pro souběžnost* (anglicky „*contracts for concurrency*“). Navržená analýza je rozdělena do dvou částí (*fáze analýzy*). **Fáze 1:** detekce *atomických sekvencí*, tj. detekce volání funkcí, které se volají atomicky. **Fáze 2:** detekce *porušení atomicity*, tj. porušení atomické sekvence získané z první fáze.

Analyzátor je implementován v jazyce *OCaml*, což je implementační jazyk nástroje Facebook Infer. Implementace se konkrétně zaměřuje na programy napsané v jazycích C/C++ s použitím zámků typu *PThread*.

Funkčnost analyzátoru byla úspěšně ověřena na menších *ručně vytvořených* programech. Navíc byl analyzátor experimentálně vyhodnocen na veřejně dostupných testovacích programech odvozených od *nízkoúrovňových* programů používaných v praxi. Analýza byla provedena na 9 nízkoúrovňových programech, které obsahují několik tisíc řádků kódu. Bylo zjištěno, že Atomer je schopný analyzovat i takhle rozsáhlé programy z praxe, ale v těchto případech je nesprávně hlášeno mnoho chyb. Každopádně výsledek této analýzy může být použit jako vstup pro *dynamickou analýzu*, která může být schopna zjistit, jestli tato nahlášená porušení atomicity jsou skutečné chyby.

Atomer prokázal svůj potenciál pro další vylepšování. Budoucí práce se bude zaměřovat především na zlepšování přesnosti použitých metod například tak, že se budou uvažovat *vnořené zámký, různé instance použitých zámků, parametry funkcí* atd. Budoucí práce se bude dále zaměřovat na zlepšování *škálovatelnosti*, protože Atomer není schopen analyzovat rozsáhlejší a komplexnější programy. Dále by bylo zajímavé rozšířit analyzátor o *další typy zámků* pro synchronizaci souběžných vláken/procesů a otestovat analýzu na dalších v praxi používaných programech.

Vývoj analyzátoru byl diskutován s vývojáři nástroje Facebook Infer a je součástí projektu Aquas (H2020 ECSEL). Zdrojové kódy implementovaného analyzátoru jsou volně dostupné v repositáři na serveru GitHub. Předběžné výsledky práce byly publikovány a prezentovány v článku na studentské konferenci *Excel@FIT*, kde tento článek vyhrál cenu ve dvou kategoriích.

Static Analysis Using Facebook Infer to Find Atomicity Violations

Declaration

Hereby I declare that this bachelor's thesis was prepared as an original author's work under the supervision of professor Tomáš Vojnar. All the relevant information sources, which were used during the preparation of this thesis, are properly cited and included in the list of references.

.....

Dominik Harmim

May 19, 2019

Acknowledgements

I would like to thank my supervisor Tomáš Vojnar. Further, I would like to thank Tomáš Fiedor for providing supplementary information and for his assistance. I would also like to thank my colleagues Vladimír Marcin and Ondřej Pavela for helpful discussions about my thesis. Furthermore, I would like to thank Nikos Gorogiannis and Sam Blackshear from Infer team at Facebook for useful discussions about the development of my analyser. Lastly, I thank for the support received from H2020 ECSEL project Aquas.

Contents

1	Introduction	2
2	Preliminaries	4
2.1	Static Analysis by Abstract Interpretation	4
2.2	Facebook Infer–Static Analysis Framework	8
2.3	Contracts for Concurrency	11
3	Atomicity Violations Detector	15
3.1	Related Work	15
3.2	Analysis and Design	16
4	Implementation and Evaluation	22
4.1	Implementation of Detection of Atomic Sequences	24
4.2	Implementation of the Detection of Atomicity Violations	28
4.3	Experimental Verification and Evaluation	32
5	Conclusion	34
	Bibliography	35
	List of Appendices	38
A	Experimental Verification Results	39
A.1	Detection of Atomic Sequences	39
A.2	Detection of Atomicity Violations	42
B	Contents of Attached Memory Media	45
C	Installation and User Manual	46
C.1	Installation Manual	46
C.2	User Manual	46

Chapter 1

Introduction

Bugs are an integral part of computer programs ever since the inception of the programming discipline. Unfortunately, they are often hidden in unexpected places, and they can lead to unexpected behaviour which may cause significant damage. Nowadays there are many possible ways of catching bugs in the development process. *Dynamic analysers* or tools for automated testing are often used. These methods are satisfactory in many cases, nevertheless, they can still leave too many bugs undetected, because they are able to analyse only certain program flows, dependent on its input data. An alternative solution is *static analysis* that has its own shortcomings as well. The main issue is the *scalability* on extensive codebases and considerable high rate of incorrectly reported errors (so-called *false positives* or *false alarms*).

Recently, Facebook introduced *Facebook Infer*—a tool for creating *highly scalable, compositional, incremental*, and *interprocedural* static analysers. Facebook Infer has grown considerably in its possibilities, but it is still under active development by many teams across the globe. It is employed every day not only in Facebook itself, but also in other companies, such as Spotify, Uber, Mozilla, or Amazon. Currently, Facebook Infer provides several analysers implemented as modules in the whole framework. These analysers check for various types of bugs, such as buffer overflows, thread-safety, null-dereferencing, or memory leaks. But most importantly Facebook Infer is a framework for building new analysers quickly and easily. Unfortunately, the current version of Facebook Infer still lacks better support for *concurrency* bugs. While it provides a fairly advanced *data race* analyser, it is limited to Java programs only and fails for C programs, which require more through manipulation with locks.

In *concurrent programs*, there are often *atomicity requirements* for execution of specific sequences of instructions. Violating these requirements may cause many kinds of problems, such as unexpected behaviour, exceptions, segmentation faults, or other failures. *Atomicity violations* are usually not verified by compilers, unlike syntactic or some sorts of semantic rules. Moreover, atomicity requirements, in most cases, are not even documented at all. So in the end, programmers themselves must abide by these requirements and usually lack any tool support. And in general, it is difficult to avoid errors in *atomicity-dependent programs*, especially in large projects, and even harder and time-consuming is finding and fixing them.

In this thesis, there is proposed the *Atomer*—the static analyser for finding atomicity violations—which is implemented as an extension for Facebook Infer. In particular, the

concentration is put on an *atomic execution of sequences of function calls*, which is often required, e.g., when using certain library calls. The implementation particularly targets C/C++ programs that use *PThread* locks.

The development of Atomer has been discussed with developers of Facebook Infer, and it is a part of the H2020 ECSEL project Aquas. Parts of this thesis and preliminary results are taken from [13], which was written in collaboration with Vladimír Marcin and Ondřej Pavela.

The rest of the paper is organised as follows. In Chapter 2, there are described all the topics which are necessary to understand before reading the rest of the paper. In particular, Section 2.1 deals with *static analysis* based on *abstract interpretation*. Facebook Infer, which uses abstract interpretation, is described in Section 2.2. And in Section 2.3, there is described the concept of *contracts for concurrency*. A proposal of a static analyser for detection of *atomicity violations*, based on this concept, is described in Chapter 3 together with a description of existing analysers of a similar kind. An implementation of the analyser and its experimental verification and evaluation is presented in Chapter 4. Finally, Chapter 5 concludes the paper. In addition, there are three appendices. Appendix A shows experimental verification results. Appendix B lists contents of attached memory media and Appendix C serves as an installation and user manual.

Chapter 2

Preliminaries

This chapter explains the theoretical background on which stands the thesis. It also explains and describes the existing tools used in the thesis. Lastly, the chapter deals with principles which this thesis got inspired by.

The aim of this thesis is to propose a *static analyser* and implement it in *Facebook Infer*. So, in Section 2.1, there is a brief explanation of *static analysis* itself, and then an explanation of *abstract interpretation* that is used in Facebook Infer. Facebook Infer, its principles and features illustrate Section 2.2. A proposal of a solution is based on the concept of *contracts for concurrency*, which is discussed and defined in Section 2.3.

2.1 Static Analysis by Abstract Interpretation

According to [19], *static analysis* of programs is reasoning about the behaviour of computer programs without actually executing them. It has been used since the 1970s for optimising compilers for generating effective code. More recently, it has proven valuable also for automatic error detection, verification tools and it is used in other tools that can help programmers. Intuitively, a static program analyser is a program that reasons about the behaviour of other programs, in other words, a static program analyser checks if the *program semantics* of a given program fulfils the given *specification*, as illustrates Figure 2.1 [8]. Nowadays, static analysis is one of the fundamental concepts of *formal verification*. It aims to automatically answer questions about a given program such as e.g. [19]:

- **Are certain operations executed *atomically*?**
- Does the program terminate on every input?
- Can the program *deadlock*?
- Does there exist an input that leads to a *null-pointer dereference*, a *division-by-zero*, or an *arithmetic overflow*?
- Are all variable initialised before they are used?
- Are arrays always accessed within their bound?

- Does the program contain *dead code*?
- Are all resources correctly released after their last use?




Figure 2.1: Static program analysis [8]

It is well-known that testing, i.e., executing programs with some input data and examining the output, may expose errors, but it can not prove their absence. (It was also famously stated by Edsger W. Dijkstra: “*Program testing can be used to show the presence of bugs, but never to show their absence!*”.) However, static program analysis can prove their absence — with some *approximation* — it can check *all possible executions* of the programs and provide guarantees about their properties. Another advantage of static analysis is that the analysis can be performed during the development process, so the program does not have to be executable yet and it already can be analysed. The significant issue is how to ensure high precision and *scalability* to be useful in practice. The biggest disadvantage is that static analysis can produce many *false alarms*¹, but it is often resolved by accepting *unsoundness*².

Various forms of static analysis of programs have been invented, for instance [24]: bug pattern searching, data-flow analysis, constraint-based analysis, type analysis, symbolic execution. And one of the essential concept — *abstract interpretation* — is detailed in Section 2.1.1.

There exist numerous tools for static analysis (often proprietary and difficult to openly evaluate or extend), e.g.: Coverity, Klockwork, CodeSonar, Loopus, phpstan, or *Facebook Infer* (described in Section 2.2).

2.1.1 Abstract Interpretation

This section explains and defines the basics of *abstract interpretation*. The description is based on [8], [9], [6], [7], [15], [16], [10], [20], [19], [25]. In these bibliographies, there also can be found more detailed, more formal, and a more theoretical explanation.

The abstract interpretation was introduced and formalised by a French computer scientist Patrick Cousot and his wife Radhia Cousot in the year 1977 at POPL³ [9]. It is a generic *framework* for static analyses. It is possible to create particular analyses by providing specific components (described later) to the framework. The analysis is guaranteed to be *sound* if certain properties of the components are met. [15], [16]

¹**False alarms** – incorrectly reported an error. Also called *false positives*.

²**Soundness** – if a verification method claims that a system is correct according to a given specification, it is truly correct. [24]

³**POPL** – symposium on Principles of Programming Languages.

In general, in the set theory, which is independent on an application setting, abstract interpretation is considered theory for *approximating* sets and set operations. A more restricted formulation of abstract interpretation is to interpret it as a theory of approximation of the behaviour of the *formal semantics* of programs. Those behaviours may be characterised by *fixpoints* (defined below), that is why a primary part of the theory provides efficient techniques for *fixpoint approximation* [20]. So, for a standard semantics, abstract interpretation is used to derive the approximate abstract semantics over an *abstract domain* (defined below), in order to check a given *program specification* using analysis of the abstract semantics. [8]

Patrick Cousot intuitively and informally illustrates abstract interpretation in [6] as follows. Figure 2.2a shows the *concrete semantics* of a program by a set of curves, which represents the set of all possible executions of the program in all possible execution environments. Each curve shows the evolution of the vector $x(t)$ of input values, state, and output values of the program as a function of the time t . *Forbidden zones* on this figure represent a set of erroneous states of the program execution. Proving, that the intersection of the concrete semantics of the program with the forbidden zone is empty, is undecidable because the program concrete semantics is not computable. As demonstrates Figure 2.2b, abstract interpretation deals with an *abstract semantics*, i.e., the *superset* of the concrete program semantics. The abstract semantics includes all possible executions. That implies that if the abstract semantics is safe (i.e. does not intersect the forbidden zone), concrete semantics is safe as well. However, the *over-approximation* of the possible program executions causes that inexistent program executions are considered, that may lead to *false alarms*. It is the case when the abstract semantics intersects the forbidden zone, whereas the concrete semantics does not intersect it.




Figure 2.2: Abstract interpretation demonstration [6]. Horizontal axes: time t . Vertical axes: vector $x(t)$ of input values of programs

Components of Abstract Interpretation

In accordance with [15], [16], basic components of abstract interpretation are as follows:

- **Abstract Domain** [7]

- An abstraction of the *concrete semantics* in the form of *abstract properties*⁴ and *abstract operations*⁵. [8]
- Sets of program states at certain locations are represented using *abstract states*.
- **Abstract Transformers**
 - There is a *transform function* for each program operation (instruction) that represents the impact of the operation executed on an abstract state.
- **Join Operator** \circ
 - Joins abstract states from individual program branches into a single one.
- **Widening Operator** ∇ [20], [10], [15]
 - Enforces termination of the abstract interpretation.
 - It is used to approximate the *least fixed points* (it is performed on a sequence of abstract states at a certain location).
 - The later in the analysis is this operator used, the more accurate is the result (but the analysis takes more time).
- **Narrowing Operator** Δ [20], [10], [15]
 - Encapsulates a termination criterion.
 - Using this operator, the approximation can be refined, i.e., it may be used to refine the result of widening.
 - This operator is used when a *fixpoint* is approximated using widening.

Fixpoints and Fixpoint Approximation

Definition 2.1.1. In [25], there is a *fixpoint* defined as:

- let (A, \leq_A) be a *lattice* [25],
- an element $a \in A$ is a **fixpoint** of a function $f : A \rightarrow A$ if and only if $f(a) = a$.

Computation of the *most precise abstract fixpoint* is not generally guaranteed to terminate in certain cases, such as loops. The solution is to approximate the fixpoint using *widening* (over-approximation of a fixpoint) and *narrowing* (improves an approximation of a fixpoint) [15], [16]. Most program properties can be represented as fixpoints. This reduces program analysis to the fixpoint approximation [7]. Further information about fixpoint approximation can be found in [20], [10].

⁴**Abstract properties** approximating *concrete properties behaviours*.

⁵**Abstract operations** include abstractions of the *concrete approximation*, an approximation of the *concrete fixpoint transform function*, etc.

Formal Definition of Abstract Interpretation

Definition 2.1.2. According to [9], [15], **abstract interpretation** I of a program P with the instruction set S is a tuple

$$I = (Q, \circ, \sqsubseteq, \top, \perp, \tau)$$

where

- Q is the *abstract domain* (domain of *abstract states*),
- $\circ : Q \times Q \rightarrow Q$ is the *join operator* for accumulation of abstract states,
- $(\sqsubseteq) \subseteq Q \times Q$ is *ordering* defined as $x \sqsubseteq y \Leftrightarrow x \circ y = y$ in (Q, \circ, \top) ,
- $\top \in Q$ is a *supremum* of Q ,
- $\perp \in Q$ is an *infimum* of Q ,
- $\tau : S \times Q \rightarrow Q$ defines the *abstract transformers* for specific instructions,
- (Q, \circ, \top) is a *complete semilattice* [25], [15].

Using so-called *Galois connections* ([20], [10], [15], [7]) can be guaranteed the *soundness* of abstract interpretation.

2.2 Facebook Infer – Static Analysis Framework

This section describes the principles and features of *Facebook Infer*. The description is based on information provided on Facebook Infer website⁶ and in [2], [16]. Parts of this section are taken from [13].

Facebook Infer is an open-source⁷ static analysis *framework*, which is able to discover various kinds of software bugs of a given program, and the stress is put on *scalability*. Elementary essence of this framework shows Figure 2.3, below is a more detailed explanation of its architecture. Facebook Infer itself is implemented in *OCaml*⁸ – *functional* programming language, also supporting *imperative* and *object-oriented* paradigms. Further details about OCaml can be found in [18] or in official documentation⁹, tutorials¹⁰. Facebook Infer was originally a standalone tool focused on *sound verification* of the absence of *memory safety violations*, which has made its breakthrough thanks to a powerful paper [5].

Facebook Infer is able to analyse programs written in several languages. In particular, it supports languages C, C++, Java, and Objective-C. Moreover, it is possible to extend Facebook Infer’s *frontend* for supporting another languages. Currently, Facebook Infer contains many analyses focusing on amount sorts of bugs, e.g., *Inferbo* (buffer overruns) [26]; *RacerD* (data races) [3], [4], [12]; and other analyses checks for buffer overflows, thread-safety, null-dereferencing, memory leaks, resource leaks, etc.

⁶Facebook Infer website – <https://fbinfer.com>.

⁷Open-source repository of Facebook Infer on GitHub – <https://github.com/facebook/infer>.

⁸OCaml website – <https://ocaml.org>.

⁹OCaml documentation – <http://caml.inria.fr/pub/docs/manual-ocaml>.

¹⁰OCaml tutorials – <https://ocaml.org/learn/tutorials>.

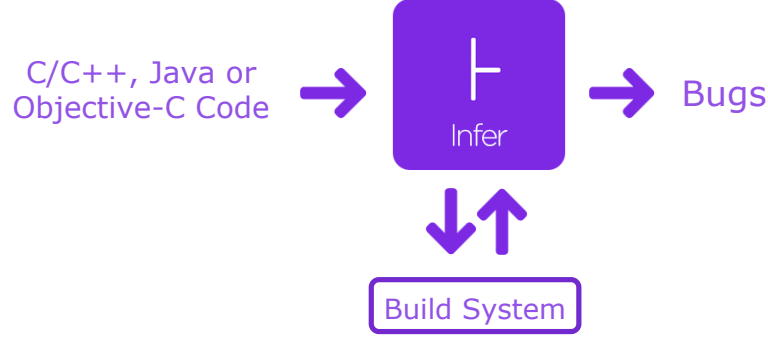


Figure 2.3: Static analysis in Facebook Infer (<http://www.codeandyou.com/2015/06/infer-static-analyzer-for-java-c-and.html>)

2.2.1 Abstract Interpretation in Facebook Infer

Facebook Infer is a general framework for static analysis of programs, it is based on *abstract interpretation*, see Section 2.1.1. It aims to find bugs rather than formal verification. It can be used to quickly develop new sorts of *compositional* and *incremental* analysers (*intraprocedural* or *interprocedural* [20]) based on the concept of function *summaries*. In general, a *summary* is a representation of *preconditions* and *postconditions* of a function. However, in practice, a summary is a custom data structure that may be used for storing any information resulting from the analysis of single functions. Facebook Infer generally does not work out the summaries in the course of the analysis along the *Control Flow Graph (CFG)*¹¹ as it is done in classical analyses based on the concepts from [21], [22]. Instead, Facebook Infer performs the analysis of a program *function-by-function along the call tree*, starting from its leafs (demonstrated later). Therefore a function is analysed and a summary is computed without knowledge of the call context. Since summaries worked out in different contexts are equal, this principle makes the analysis more scalable, but it can lead to a loss of accuracy. Then, the summary of a function is used at all of its call sites. In order to create new intraprocedural analyser in Facebook Infer, it is needed to define (listed items are described in more detail in Section 2.1.1):

1. The *abstract domain* Q , i.e., a type of an *abstract state*.
2. Operator \sqsubseteq , i.e., *ordering* of abstract states.
3. The *join* operator \circ , i.e., the way of joining two abstract states.
4. The *widening* operator ∇ , i.e., the way how to enforce termination of the abstract interpretation of iteration.
5. *Transfer functions* τ , i.e., a transformer that takes an abstract state as an input and produces an abstract state as an output.

And in order to create an interprocedural analyser, it is required to additionally define:

1. A type of function summaries.

¹¹A **control flow graph (CFG)** is a directed graph in which the nodes represent basic blocks and the edges represent control flow paths. [1]

2. The logic for using summaries in transfer functions, and the logic for transforming an intraprocedural abstract state to a summary.

The next important feature improving the scalability is *incrementality* of the analysis, it allows to analyse separate code changes only, instead of analysing the whole codebase. It is more suitable for extensive and variable projects, where ordinary analysis is not feasible. The incrementality is based on *re-using summaries* of functions for which there is no change in them neither in the functions transitively invoked from them.

The Architecture of the Abstract Interpretation Framework in Facebook Infer

The architecture of the abstract interpretation framework of Facebook Infer (**Infer.AI**) may be split into three major parts, as demonstrates Figure 2.4: a *frontend*, an *analysis scheduler* (and a *results database*), and a set of *analyser plugins*.

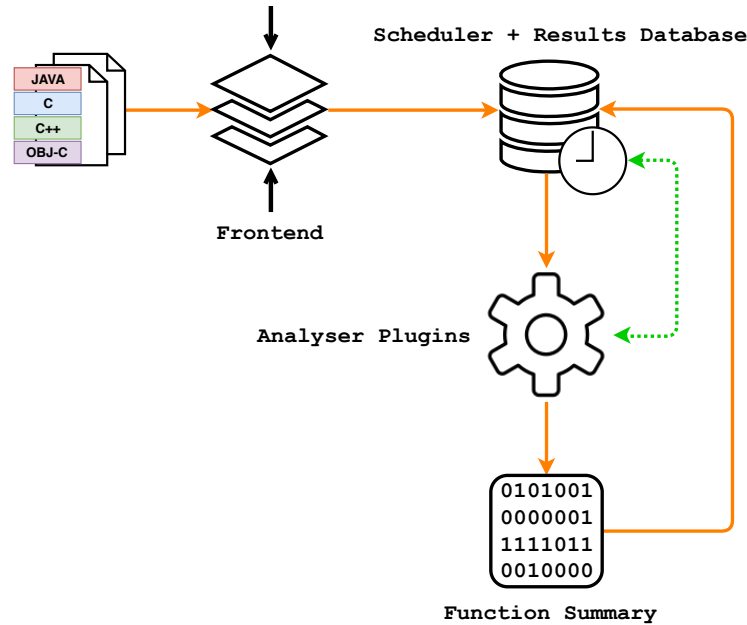


Figure 2.4: The architecture of Facebook Infer’s abstract interpretation framework [2], [16]

The frontend compiles input programs into the *Smallfoot Intermediate Language* (SIL) and represents them as the CFG. There is a separate CFG representation for each analysed function. Nodes of this CFG are formed as instructions of SIL. SIL language consists of following underlying instructions:

1. **LOAD** – reading into a temporary variable.
2. **STORE** – writing to a program variable, a field of a structure, or an array.
3. **PRUNE e** (often called **ASSUME**) – a condition **e**.
4. **CALL** – a function call.

The frontend allows one to propose *language-independent* analyses (to a certain extent) because it supports input programs to be written in multiple languages.

The next part of the architecture is the scheduler, which defines the order of the analysis of single functions according to the appropriate *call graph*¹². The scheduler also checks if it is possible to analyse some functions simultaneously, which allows Facebook Infer to run the analysis in parallel.

Example 2.2.1. For demonstrating the order of the analysis in Facebook Infer and its incrementality, assume a call graph in Figure 2.5. At first, leaf functions **F5** and **F6** are analysed. Further, the analysis goes on towards the root of the call graph—**F_{MAIN}**, while takes into consideration the dependencies denoted by the edges. This order ensures that a summary is available once a nested function call is abstractly interpreted within the analysis. When there is a subsequent code change, only directly changed functions and all the functions up the call path are re-analysed. For instance, if there is a change of source code of function **F4**, Facebook Infer triggers re-analysis of functions **F4**, **F2**, and **F_{MAIN}** only.




Figure 2.5: A call graph for an illustration of Facebook Infer’s analysis process [2], [13], [16]

The last part of the architecture consists of a set of analyser plugins. Each plugin performs the analysis by interpretation of SIL instructions. Result of the analysis of each function (function summary) is stored to the results database. Interpretation of SIL instructions (*commands*) is done using an *abstract interpreter* (also called a *control interpreter*) and *transfer functions* (also called a *command interpreter*). The transfer functions take an actual *abstract state* of an analysed function as an input, and by applying the interpreting command produce a new abstract state. Then, the abstract interpreter interprets the command in *abstract domain* according to the CFG. This workflow is simplified in Figure 2.6.

2.3 Contracts for Concurrency

This section introduces and defines the concept of *contracts for concurrency* described in [23], [11]. Parts of this section are taken from [13]. Listings in this section are pieces of programs written in ANSI C¹³.

Respecting the *protocol* of a software module—delineates which *sequences of functions* are legal to invoke—is one of the requirements for the correct behaviour of the module. For example, a module that deals with file system typically requires that a programmer using this module should call function **open** at first, followed by an optional number of functions **read** and **write**, and at last, call function **close**. A program utilising such a module that does not follow this protocol is erroneous. The methodology of *design by contract* (described in [17]) requires programs to meet such well-defined behaviours. [23]

¹²A **call graph** is *directed graph* describing call dependencies among functions.

¹³**ANSI C**—standard for the C programming language published by the *ANSI* (American National Standards Institute).




Figure 2.6: Facebook Infer’s abstract interpretation process [2], [16]

In *concurrent programs*, contracts for concurrency allow one to specify *sequences of functions* that are needed to be *executed atomically*, in order to avoid *atomicity violations*. Such contracts may be manually specified by a developer or it may be automatically generated by a program (analyser). These contracts can be used to verify the correctness of programs as well as they can serve as helpful documentation. A program is safe from atomicity violations if the program follows the contract and the contract is well-defined and complete.

Section 2.3.1 defines the notion of *basic contracts for concurrency*. Further, Section 2.3.2 defines contracts extended to consider the *data flow* between functions (i.e., a sequence of function calls must be atomic only if they handle the same data). Above that, paper [11] extends the idea of basic contracts with *spoilers* (i.e., extending by *contextual information*).

2.3.1 Basic Contracts

Definition 2.3.1. In [11], [23], a *basic contract* is formally defined as follows. Let $\Sigma_{\mathbb{M}}$ be a set of all function names of a software module. A contract is a set \mathbb{R} of *clauses* where each clause $\varrho \in \mathbb{R}$ is a *star-free regular expression*¹⁴ over $\Sigma_{\mathbb{M}}$. A *contract violation* occurs if any of the sequences expressed by the contract clauses are interleaved with the execution of functions from $\Sigma_{\mathbb{M}}$, in other words, each sequence specified by any clause ϱ must be executed atomically, otherwise, there is a violation of the contract. The number of sequences defined by a contract is finite since the contract is the union of *star-free languages*.

Example 2.3.1. Consider the following example from [11], [23]. There is a module with the implementation of a resizable array with the listed functions:

¹⁴**Star-free regular expressions** are regular expressions using only the *concatenation operators* and the *alternative operators* ($()$), without the *Kleene star operator* ($*$).

```

f1: void add(char *array, char element)
f2: bool contains(char *array, char element)
f3: int index_of(char *array, char element)
f4: char get(char *array, int index)
f5: void set(char *array, int index, char element)
f6: void remove(char *array, int index)
f7: int size(char *array)

```

The module's contract contains the following clauses:

(ϱ_1) contains index_of

The execution of `contains` followed by the execution of `index_of` should be atomic. Otherwise, the program may fail to get the index, because after verification of the presence of an element in an array, it can be concurrently, e.g., removed.

(ϱ_2) index_of (get | set | remove)

The execution of `index_of` followed by the execution of `get`, `set`, or `remove` should be atomic. Otherwise, the received index may be outdated when it is applied to address an element, because a concurrent modification of an array may shift the position of the element.

(ϱ_3) size (get | set | remove)

The execution of `size` followed by the execution of `get`, `set`, or `remove` should be atomic. Otherwise, the size of an array may be void when accessing an array, because of a concurrent change of the array. This can be an issue since a given index is not in a valid range anymore (e.g., testing `index < size`).

(ϱ_4) add (get | index_of)

The execution of `add` followed by the execution of `get` or `index_of` should be atomic. Otherwise, the added element does not have to longer exist or its position in an array can be changed, when the program attempts to obtain information about it.

The above definition of contracts for concurrency is quite limited in some circumstances and can consider valid concurrent programs as erroneous (reports *false alarms*). Hence, in Section 2.3.2, there is defined an extension of contracts for concurrency with *parameters*, which takes into consideration the data flow within function calls. And in [11], [23], there is defined another extension with *spoilers*, which considering contextual information of function calls.

2.3.2 Contracts with Parameters

Example 2.3.2. Consider the following example from [11], [23], as demonstrates Listing 2.1. There is a function `replace` that replaces item `a` in an array by item `b`. Implementation of this function comprises two atomicity violations:

- (i) when `index_of` is invoked, item `a` does not need to be in the array anymore;
- (ii) the acquired index can be obsolete when `set` is invoked.

A basic contract defined in Section 2.3.1 could cover this scenario by clause ϱ_5 :

(ϱ_5) contains index_of set

Nevertheless, it is too restrictive because it is required to be executed atomically only if `contains` and `index_of` have the same arguments `array` and `element`, `index_of` and `set` have the same argument `array`, and the returned value of `index_of` is used as the argument `index` of function `set`.

```

1 void replace(char *array, char a, char b)
2 {
3     if (contains(array, a))
4     {
5         int index = index_of(array, a);
6         set(array, index, b);
7     }
8 }

```

Listing 2.1: An example of an atomicity violation with data dependencies [11]

In order to respect function call *parameters* and *return values* of functions in contracts, the basic contracts are further extended by dependencies among functions in [11], [23] as follows. Function call parameters and return values are expressed as *meta-variables*. Further, if a contract should be required exclusively if the same object emerges as an argument or as the return value of multiple calls in a given call sequence, it may be denoted by using the same meta-variable at the position of all these occurrences of parameters and return values.

Clause ϱ_5 can be extended as follows (repeated application of meta-variables `X/Y/Z` requiring the same objects `o1/o2/o3` to be used at the positions of `X/Y/Z`):

(ϱ'_5) contains(X,Y) Z=index_of(X,Y) set(X,Z,_)

The underscore indicates a *free meta-variable* that does not restrict the contract clause.

With the extension described above, it is possible to extend the contract from Section 2.3.1 as follows:

(ϱ'_1) contains(X,Y) index_of(X,Y)

(ϱ'_2) Y=index_of(X,_) (get(X,Y) | set(X,Y,_) | remove(X,Y))

Chapter 3

Atomicity Violations Detector

This chapter describes a proposal of a static program analyser for detection of *atomicity violations*. The proposed analyser — **Atomer** — has been proposed as an extension for *Facebook Infer*, introduced in Section 2.2. In particular, the proposal concentrates on an *atomic execution of sequences of function calls*, which is often required. The proposed principle is based on the assumption that sequences executed *once atomically* should probably be executed *always atomically*. The chapter also assesses already existing solutions in this area.

At first, Section 3.1 deals with existing approaches and tools for finding atomicity violations, their advantages, disadvantages, features, availability, and so on. Then, the proposal itself is introduced in Section 3.2. Parts of this chapter are taken from [13]. Listings in this chapter are pieces of exemplary programs written in ANSI C (assume *PThread* locks and the existence of an initialised global variable `lock` of a type `pthread_mutex_t`).

3.1 Related Work

The proposed solution is slightly inspired by ideas from [11], [23]. In these papers, there is described a proposal and implementation of a *static validation* for finding *atomicity violations*, which is based on *grammars* and *parsing trees*. In paper [11], there is also described and implemented a dynamic approach to this validation. The authors of [11] and [23] implemented a stand-alone prototype tool¹ for analysing programs written in Java. It led to some promising experimental results but the *scalability* of the tool was still limited. Moreover, the tool from [11] and [23] is no more developed. That is why was made the decision to get inspired by [11] and [23] and reimplement the analysis in *Facebook Infer* redesigning it in accordance with the principles of Facebook Infer (described in Section 2.2), which should make it more scalable. In the end, due to adapting the analysis for the context of Facebook Infer, implementation of the analysis within this thesis is significantly different from [11] and [23], as it is presented in Chapter 4. Furthermore, unlike [11] and [23], the implementation aims at programs written in C/C++ languages using *POSIX Thread* (*PThread*) locks for a *synchronisation of concurrent threads*.

¹**Gluon** — a tool for static verification of *contracts for concurrency* (see Section 2.3) in Java programs — <https://github.com/trxsys/gluon>.

In Facebook Infer, there is already implemented analysis called *Lock Consistency Violation*². It is part of *RacerD* [3], [4], [12]. This analysis finds atomicity violations for writes/reads on single variables that are required to be executed atomically. Atomer is different, it finds atomicity violations for *sequences of functions* that are required to be executed atomically, i.e., it checks whether *contracts for concurrency* (see Section 2.3) hold.

3.2 Analysis and Design

The proposal of the analyser is based on the concept of *contracts for concurrency* described in Section 2.3. In particular, the proposal considers only *basic contracts* described in Section 2.3.1. Parameters of functions and their return values—expressed in *contracts with parameters* (see Section 2.3.2)—are not taken into consideration.

In general, basic contracts for concurrency allow one to define *sequences of functions* that are required to be *executed atomically*, as it is explained in more detail in Section 2.3. Atomer is able to automatically derive candidates for such contracts, and then to verify whether the contracts are fulfilled. Both of these operations are done statically. The proposed analysis is divided into two parts (*phases of the analysis*):

Phase 1: Detection of *atomic sequences*, which is described in Section 3.2.1.

Phase 2: Detection of *atomicity violations* (violations of the atomic sequences), which is described in Section 3.2.2.

These phases of the analysis and its workflow illustrate Figure 3.1.

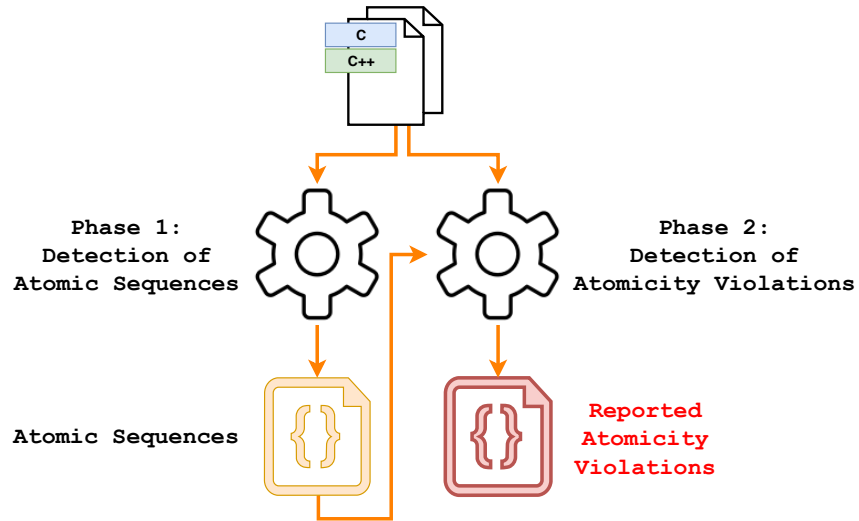


Figure 3.1: Phases of the proposed analyser

²**Lock Consistency Violation**—atomicity violations analysis in Facebook Infer—https://fbinfer.com/docs/checkers-bug-types.html#LOCK_CONSISTENCY_VIOLATION.

3.2.1 Phase 1: Detection of Atomic Sequences

Before the detection of *atomicity violations* (Section 3.2.2) may begin, it is required to have contracts introduced in Section 2.3. **Phase 1** of Atomer is able to produce such contracts, i.e., it detects *sequences of functions* that should be *executed atomically*. Intuitively, the detection is based on looking for sequences of functions that are executed atomically on some path through a program. The assumption is that if it is once needed to execute a sequence atomically, it should probably be always executed atomically.

The detection of sequences of calls to be executed atomically is based on analysing all paths through the CFG of a function and generating all pairs **(A, B)** of sets of function calls such that: **A** is a *reduced sequence* of function calls that appear between the beginning of the function being analysed and the first lock or between an unlock and a subsequent lock (or between an unlock and the end of the function being analysed), and **B** is a reduced sequence of function calls that follow the calls from **A** and that appear between a lock and an unlock (or between a lock and the end of the function being analysed). Here, by a reduced sequence, it is meant a sequence in which the first appearance of each function is recorded only. The reason is to ensure *finiteness* of the sequences and of the analysis. The *summary* of a function then consists of:

- (i) the set of all the **B** sequences and
- (ii) the set of *concatenations* of all the **A** and **B** sequences with the removal of duplicate function calls.

The latter is recorded for the purpose of analysing functions higher in the *call hierarchy* since locks/unlocks can appear in such a *higher-level function*.

Example 3.2.1. For instance, the analysis of the function **g** from Listing 3.1 produces the following sequences:

$$\overbrace{\text{f1 f1}}^{\mathbf{A}_1} \overbrace{(\text{f1 f1 f2})}^{\mathbf{B}_1} \mid \overbrace{\text{f1 f1}}^{\mathbf{A}_2} \overbrace{(\text{f1 f3})}^{\mathbf{B}_2} \mid \overbrace{\text{f1}}^{\mathbf{A}_3} \overbrace{(\text{f1 f3 f3})}^{\mathbf{B}_3}$$

The parentheses are used to indicate an atomic sequence. The strikethrough of the functions **f1** and **f3** denotes the removal of already recorded function calls in the **A** and **B** sequences. The strikethrough of the entire sequence **f1 (f1 f3 f3)** means discarding sequences already seen before. The derived sets for the function **g** are then as follows:

- (i) $\{(\text{f1 f2}), (\text{f1 f3})\}$, i.e., **B₁** and **B₂**;
- (ii) $\{\text{f1 f2 f3}\}$, i.e., concatenation of **A₁**, **B₁**, **A₂**, and **B₂** with the removal of duplicate function calls.

Analysing Functions Using Results of the Analysis of Nested Functions

Further, it is demonstrated how the results of the analysis of *nested functions* are used during the detection of atomic sequences. The result of the analysis of a nested function is used as follows. When calling an already analysed function, one plugs all the sequences from the second component of its summary into the current **A** or **B** sequence.

```

1 void g(void)
2 {
3     f1(); f1();
4
5     pthread_mutex_lock(&lock);
6     f1(); f1(); f2();
7     pthread_mutex_unlock(&lock);
8
9     f1(); f1();
10
11    pthread_mutex_lock(&lock);
12    f1(); f3();
13    pthread_mutex_unlock(&lock);
14
15    f1();
16
17    pthread_mutex_lock(&lock);
18    f1(); f3(); f3();
19    pthread_mutex_unlock(&lock);
20 }

```

Listing 3.1: An example of a code for an illustration of the derivation of sequences of functions called atomically

Example 3.2.2. This example shows how the function **h** from Listing 3.2 would be analysed using the result of the analysis of the function **g** from Listing 3.1. So the analysis of the function **h** produces the following sequence:

$$f1 \ g \ f1 \ f2 \ f3 \ (g \ f1 \ f2 \ f3)$$

The derived sets for the function **h** are then as follows:

- (i) $\{(g \ f1 \ f2 \ f3)\}$;
- (ii) $\{f1 \ g \ f2 \ f3\}$.

```

1 void h(void)
2 {
3     f1(); g();
4
5     pthread_mutex_lock(&lock);
6     g();
7     pthread_mutex_unlock(&lock);
8 }

```

Listing 3.2: An example of a code for an illustration of the derivation of sequences of functions called atomically with a nested function call (function **g** is defined in Listing 3.1)

Cases Where Lock/Unlock Calls Are Not Paired in a Function

For treating cases where lock/unlock calls are *not paired* in a function — as demonstrates Listing 3.3 — two solutions have been proposed:

1. At the end of a function, everything is unlocked, i.e., append an unlock to the end of the function if it is necessary. Then for the function `x` from Listing 3.3, the first component of its summary (i.e., atomic sequences) would be $\{(a)\}$. Subsequently, all unlock calls not preceded by a lock are ignored. So the first component of a summary of the function `y` from Listing 3.3 would be an empty set.
2. Addition of two further items to the summaries:
 - (a) function calls with missing an unlock call,
 - (b) function calls with missing a lock call.

For the example from Listing 3.3, this would give:

- for `x`: $\{(f1)\}$,
- for `y`: $\{(f2)\}$.

The above sequences would have to be glued to the sequences captured higher in the call hierarchy. Calls of the functions `f1` and `f2` will also appear in the second component of the function summaries (i.e., the sequences of all functions called).

```
1 void x(void)
2 {
3     pthread_mutex_lock(&lock);
4     f1();
5 }
6
7 void y(void)
8 {
9     f2();
10    pthread_mutex_unlock(&lock);
11 }
12
13 void main(void)
14 {
15     x(); y();
16 }
```

Listing 3.3: An example of a code for an illustration of treating cases where lock/unlock calls are not paired in a function

In the end, the first approach of treating such cases described above has been chosen. The reason is that it is much easier for implementation. However, in future, the analysis can be improved by implementing the second approach.

Summary of Detection of Atomic Sequences and Future Work

The above detection of atomic sequences has been implemented, as it is described in Section 4.1. Furthermore, it has been successfully verified on a set of sample programs created for this purpose. The verification is presented in Sections 4.3, A.1. The derived sequences of calls assumed to execute atomically, i.e., the **B** sequences, from the summaries of all analysed functions are stored into a file, which is used during **Phase 2**, described below in Section 3.2.2. There are some possibilities for further extending and improving **Phase 1**, e.g., working with *nested locks*; distinguishing the *different locks* used (currently, it is not distinguished between the locks at all); consider *contracts for concurrency with parameters* defined in Section 2.3.2 or other extensions of contracts for concurrency discussed in Section 2.3; or extending the detection for *other types of locks* for synchronisation of concurrent threads/processes. On the other hand, to further enhance the *scalability*, it seems promising to replace working with the **A** and **B** sequence by working with sets of calls: sacrificing some precision but gaining the speed.

3.2.2 Phase 2: Detection of Atomicity Violations

In the second phase of the analysis, i.e., when *detecting violations* of the atomic sequences obtained from **Phase 1** (see Section 3.2.1), the analysis looks for pairs of functions that should be called atomically (or just for single functions if there is only one function call in an atomic sequence) while this is not the case on some path through the CFG.

Example 3.2.3. For example, assume that the result of the first phase is the following set of functions called atomically:

$$\{(\mathbf{f1\ f2\ f3}), (\mathbf{f1\ f3\ f4})\}$$

Then the analysis will look for the following pairs of functions that are not called atomically:

- **f1 f2**
- **f2 f3**
- **f1 f3**
- **f3 f4**

The analysis of functions with nested function calls and cases where lock/unlock calls are not paired in a function are handled the analogical way as it is handled in **Phase 1** described in Section 3.2.1. For detailed examples see verification experiments in Section A.2.

Example 3.2.4. For a demonstration of the detection of an atomicity violation, assume the functions **a** and **b** from Listing 3.4. The set of atomic sequences of the function **a** is $\{(\mathbf{f2\ f3})\}$. In the function **b**, an atomicity violation is detected because the functions **f2** and **f3** are not called atomically (under a lock).

```

1 void a(void)
2 {
3     f1();
4
5     pthread_mutex_lock(&lock);
6     f2(); f3();
7     pthread_mutex_unlock(&lock);
8
9     f4();
10 }
11
12 void b(void)
13 {
14     f1(); f2(); f3(); f4();
15 }

```

Listing 3.4: Example of an atomicity violation

Summary of Detection of Atomicity Violations and Future Work

As well as the first phase of the analysis, **Phase 2** has been implemented, as it is described in Section 4.2. And it has been also successfully verified on a set of sample purposeful programs. This verification is described in Sections 4.3, A.2. **Phase 2** also has the potential for further enhancing. It is possible to extend this phase with all the improvements discussed in Section 3.2.1. Moreover, it is possible to improve this phase by working with sets instead of pairs, when looking for functions that should be called atomically. The next idea is to consider atomic sequences from the first phase only if they appear in an *atomic block* more than, e.g., three times. It would strengthen certainty that this sequence should be called atomically.

Chapter 4

Implementation and Evaluation

This chapter describes the implementation of the *static analyser* proposed in Chapter 3. The analyser is implemented as an extension for *Facebook Infer* introduced in Section 2.2. The implementation is demonstrated using algorithms in *pseudocode* and using listings with codes written in *OCaml*, which is an implementation language of Facebook Infer. Section 4.1, respectively Section 4.2, then describes an implementation of the *detection of atomic sequences* defined in Section 3.2.1, respectively an implementation of the *detection of atomicity violations* defined in Section 3.2.2. Subsequently, Section 4.3 covers experimental verification and evaluation of the analyser.

The implementation of the analyser can be found publicly on GitHub¹. The implementation is done in OCaml and it is exploited its *functional* and *imperative* paradigm. Facebook Infer supports analysis of programs written in Java, C, C++, and Objective-C. However, the implementation aims at programs written in C/C++ languages using *POSIX Thread (PThread)* locks, which is a low-level mechanism for *synchronisation of concurrent threads*. So, as a lock, it is considered a function with a name `pthread_mutex_lock` and as an unlock, it is considered a function with a name `pthread_mutex_unlock`. It is also possible to run the analysis on programs written in Java or Objective-C languages but the result of the analysis would be likely wrong since these languages use a different mechanism for synchronisation.

Phase 1, i.e., the detection of atomic sequences and **Phase 2**, i.e., the detection of atomicity violations (these phases are described in Chapter 3.2) are implemented as separated analysers in Facebook Infer. The output of the first phase is the input of the second phase (as demonstrates Figure 3.1). Both of these analysers are registered as extensions of Facebook Infer in a file `infer/src/checkers/registerCheckers.ml`. These analyses run only if a particular command line argument of Facebook Infer is specified. Implementations of individual phases are stated in sections below (Section 4.1 and Section 4.2).

In order to make the analysis *interprocedural*, it is necessary to define a type of function *summaries* for each phase. The types of summaries are defined in *abstract domains* of each phase. But the summaries are accessed globally using a structure (so-called *summary payload*). Fields of the payload that refer to the summaries of analyses are defined in a file `infer/src/backend/Payloads.ml`.

¹The implementation of the analyser in a GitHub repository, which is a *fork* of the official repository of Facebook Infer, in a branch `atomicity`—<https://github.com/harmim/infer/tree/atomicity>.

For both phases of the analysis, the analyser is implemented as an *abstract interpreter* using the `LowerHil` module which transforms *SIL instructions* into *HIL instructions*. (Abstract interpretation in Facebook Infer, as well as SIL instructions, are described in Section 2.2.1.) HIL instructions just wrap SIL instructions and simplify their utilisation. For representing functions, *forward CFG with no exceptional control-flow* is used. It corresponds to the `ProcCfg.Normal` module in Facebook Infer. *Transfer functions* of both phases are implemented in the same essence, as illustrates Listing 4.1. In general, transfer functions take an *abstract state* as an input and produce an abstract state as an output for specific instructions. In this case, it modifies an abstract state when a function is called (`CALL` instruction). When the called function is a lock or an unlock, the abstract state is appropriately updated in the abstract domain of the analysis. Otherwise, the abstract state is updated with the called function and then, if the called function has already been analysed, it is used its summary to update the abstract state again.

```

1 let exec_instr astate procData _ instr =
2   match instr with
3   | Call (_, Direct procName, _, _, _) ->
4     let procNameS = Procname.to_string procName in
5
6     if is_lock procNameS then
7       Domain.update_astate_on_lock astate
8     else if is_unlock procNameS then
9       Domain.update_astate_on_unlock astate
10    else
11      (
12        let astate =
13          Domain.update_astate_on_function_call astate procNameS
14        in
15
16        match Payload.read procData.pdesc procName with
17        | Some summary ->
18          Domain.update_astate_on_function_call_with_summary astate summary
19        | None -> astate
20      )
21    | _ -> astate

```

Listing 4.1: *Transfer functions* of the analysers

The abstract domains of both phases are altogether dissimilar. But essential *operators* of the abstract domains are practically the same. Implementation of these operators is put forward in Listing 4.2, where `TSet` is a module representing a set of specific structures. The abstract state is then of a type of `TSet`. Each phase of the analysis define its own `TSet`, i.e., fields of structures in this set are different for each phase. So particular operators are defined as follows (see also Listing 4.2):

- The *ordering* operator \sqsubseteq (in Facebook Infer, it is `<=`) is defined as follows. Let `lhs` be a left-hand side of this operator and `rhs` a right-hand side of this operator. Then, `lhs <= rhs` (`lhs` is less or equal to `rhs`) if and only if `lhs` is a *subset* of `rhs`.

- The *join* operator \circ (in Facebook infer, it is `join`) is defined as a *union* of two abstract states.
- The *widening* operator ∇ (in Facebook Infer, it is `widen`) is defined as *join* of the previous and next abstract states.

```

1 let ( <= ) ~lhs:leftSide ~rhs:rightSide =
2   TSet.is_subset leftSide ~of_:rightSide
3
4 let join astate1 astate2 =
5   TSet.union astate1 astate2
6
7 let widen ~prev:prevAstate ~next:nextAstate ~num_iters:_ =
8   join prevAstate nextAstate

```

Listing 4.2: Essential *operators of abstract domains* of the analysers

4.1 Implementation of Detection of Atomic Sequences

The proposal of this phase is described in Section 3.2.1. It runs only if it is specified a command line argument `--atomic-sequences`. It detects *sequences of functions* that should be *executed atomically*. These sequences are printed into a file, it is explained in Section 4.1.2.

The main function of the analyser of this phase is `analyse_procedure`, which is shown in Listing 4.3. Facebook Infer invokes this function for every function in an analysed program. It produces a *summary* for a given function. The function computes an *abstract state* for the analysed function using the created abstract interpreter `Analyser` on an *abstract domain*. As a *precondition*, an initial abstract state `initialAstate` from the abstract domain is used. If the computation succeeds, the abstract state is appropriately updated and converted to the function summary by applying functions in the abstract domain. In the end, the *summary payload* is updated with the resulting summary.

The abstract domain of this phase is described in Section 4.1.1. It includes the definition of an abstract state, summary, and functions working with them. *Ordering* of abstract states, the *join operator*, and the *widening operator* are defined at the beginning of Chapter 4.

4.1.1 Abstract Domain of the Detection of Atomic Sequences

In this section, at first, it is described the definition of an *abstract state* of the *abstract domain* along with functions working with the abstract state. Furthermore, it is described a *summary* of functions in this phase of the analysis and corresponding functions working with the summary.

```

1 let analyse_procedure args =
2   let procData = ProcData.make_default args.proc_desc args.tenv in
3
4   match Analyser.compute_post procData ~initial:Domain.initialAstate with
5   | Some astate ->
6     let summary =
7       let astate = Domain.update_astate_at_the_end_of_function astate in
8       Domain.convert_astate_to_summary astate
9     in
10
11    Payload.update_summary summary args.summary
12  | None -> Logging.(die InternalError) "Analysis failed."

```

Listing 4.3: The analysis of a function in the analyser of **Phase 1**

Abstract State of the Domain of the Detection of Atomic Sequences

The abstract state is of a type of `TSet`. `TSet` is a module representing a **set of structures**. This structure has the following fields:

- **firstOccurrences**—it is of a type of a **list of strings**. It captures the *first occurrences* of function calls in the **A** or **B** sequences defined in Section 3.2.1. In other words, it captures the first occurrences of function calls inside or outside *atomic blocks*.
- **callSequence**—it is of a type of a **list of strings**. It is used for storing the **A** sequences followed by the **B** sequences. In other words, it stores function calls outside atomic blocks followed by function calls inside atomic blocks. For instance, `f1 f2 (f3)`.
- **finalCalls**—it is of a type of a **set of lists of strings**. It is used for storing a set of sequences of calls **callSequence**. For instance, `{f1 f2 (f3), f2 (f1 f3)}`.
- **isInLock**—it is of a type of a **boolean**. Determines whether the current state of a function is inside or outside an atomic block, i.e., it is or it is not under a lock.

The *initial abstract state* is then a set with a single empty element. An empty element is an element where **firstOccurrences** and **callSequence** are empty strings, **finalCalls** is an empty set, and **isInLock** is `false`.

According to Listings 4.1 and 4.3, there are several functions working with the abstract state. The functions are described below (all of these functions modifies all elements of the abstract state):

- **update_astate_on_function_call**—it is invoked when any function (except a lock or an unlock) is called. It captures the first occurrence of the called function.
- **update_astate_on_lock**—it is invoked when a lock is called. When the state is not under a lock, it sets a flag indicating the start of an atomic sequence. Moreover, capturing the first occurrences of function calls inside the atomic sequence begins.

- `update_astate_on_unlock`—it is invoked when an unlock is called. When the state is under a lock, it unsets a flag indicating the start of an atomic sequence. Moreover, capturing of the first occurrences of function calls followed this unlock call begins, and the last captured function calls, i.e., the last captured **A** and **B** sequences, are moved into the set of all such captured sequences within an analysed function.
- `update_astate_at_the_end_of_function`—it is invoked at the end of the analysis of a function. It moves the last captured function calls, i.e., the last captured **A** and **B** sequences, into the set of all such captured sequences within an analysed function.

Function Summary of the Domain of the Detection of Atomic Sequences

The summary is of a type of **structure**. The structure has the following fields:

- `atomicSequences`—it is of a type of a **list of lists of strings**. It is a list of all captured atomic sequences within an analysed function. For instance, `(f3) (f1 f3)`. It is used for printing the atomic sequences to a file at the end of the entire analysis.
- `allOccurrences`—it is of a type of a **list of strings**. It is a list of all called functions within an analysed function. It is used for the purpose of analysing functions higher in the *call hierarchy*.

According to Listings 4.1 and 4.3, there are some functions working with the summary. These functions are described below:

- `update_astate_on_function_call_with_summary`—it is invoked when the called function has already been analysed so that the abstract state could be updated with its summary. Therefore, occurrences of all called functions within the called function are appended to the first occurrences of an analysed function. It is demonstrated on Algorithm 4.1.
- `convert_astate_to_summary`—it is invoked at the end of the analysis of an analysed function. It transforms the abstract state of a given function to the summary. In particular, it derives all atomic sequences and all called functions within an analysed function from the abstract state, as demonstrates Algorithm 4.2.

Algorithm 4.1: Updating the abstract state with the summary of a called function

```
1 def update_astate_on_function_call_with_summary(astate, sum):
2   if sum.allOccurrences ≠ [] then
3     for e ∈ astate do
4       for o ∈ sum.allOccurrences do
5         | e.firstOccurrences ← AddUniq(e.firstOccurrences, o);
6       end
8     end
9   return astate;
10 end
```

Algorithm 4.2: Converting the abstract state to the function summary

```
1 def convert_astate_to_summary(astate):
2   atomicSeq ← [];
3   allOccur ← [];
4   for e ∈ astate do
5     for c ∈ e.finalCalls do
6       | atomicSeq ← AddUniq(atomicSeq, GetAtomicSeq(c));
7       | allOccur ← AddUniq(allOccur, GetAllCalls(c));
8     end
9   end
10  return {atomicSeq, allOccur};
11 end
```

4.1.2 Output of the Detection of Atomic Sequences

The output of **Phase 1** are sequences of functions that should be executed atomically for each analysed function in a program. These sequences are derived from summaries of all analysed functions. At the end of the entire analysis, the sequences are printed into a file `infer-atomicity-out/atomic-sequences` in the following format. Each line of the file contains a list of the detected atomic sequences within a particular function. It starts by a function name of an analysed function followed by a colon and whitespace. Then, there are listed atomic sequences (function names separated by whitespace) separated by whitespace. Example of the output:

```
functionA:_(f1_f2)_(f3_f1)
functionB:_
functionC:_(f3_f4)_(f6)
```

The principle of the derivation of the atomic sequences and their printing is demonstrated on Algorithm 4.3. The atomic sequences are then further processed in the second phase of the analysis, see Section 4.2.

Algorithm 4.3: Printing atomic sequences from summaries of all analysed functions

Input: A set F of all analysed functions

```
1 for  $f \in F$  do
2   printf('%s:␣', GetFunName( $f$ ));
3    $S \leftarrow \text{ReadSummary}(f)$ ;
4   for  $q \in S.\text{atomicSequences}$  do
5     printf('(%s)␣', SeqToString( $q$ ));
6   end
7 end
```

4.2 Implementation of the Detection of Atomicity Violations

The proposal of this phase is described in Section 3.2.2. It runs only if it is specified a command line argument `--atomicity-violations`. It detects *atomicity violations*, i.e., violations of the *atomic sequences* obtained from **Phase 1**. The atomic sequences are read from the file `infer-atomicity-out/atomic-sequences` (see Section 4.1.2). If this file does not exist, i.e., the previous phase of the analysis has not run yet, this phase will fail.

As well as it is with the first phase of the analysis, the main function of the analyser of this phase is `analyse_procedure`, which is shown in Listing 4.4. Facebook Infer invokes this function for every single function in an analysed program. And it produces a summary for a given function. This function, at first, initialises an abstract domain of this phase, and then it computes an abstract state for an analysed function using the created abstract interpreter **Analyser** upon the abstract domain. As a *precondition*, an initial abstract state `initialAstate` from the abstract domain is used. If the computation succeeds, the abstract state is converted to the function summary by application of functions in the abstract domain. Further, atomicity violations within the analysed function are reported based on the abstract state. This reporting is in more detail described in Section 4.2.2. In the end, the *summary payload* is updated with the resulting summary.

```
1 let analyse_procedure args =
2   Domain.initialise true;
3   let procData = ProcData.make_default args.proc_desc args.tenv in
4
5   match Analyser.compute_post procData ~initial:Domain.initialAstate with
6   | Some astate ->
7     let summary = Domain.convert_astate_to_summary astate in
8
9     Domain.report_atomicity_violations astate ( fun loc msg ->
10       Reporting.log_error
11         args.summary ~loc:loc IssueType.atomicity_violation msg );
12
13     Payload.update_summary summary args.summary
14   | None -> Logging.(die InternalError) "Analysis failed."
```

Listing 4.4: The analysis of a function in the analyser of **Phase 2**

The abstract domain of this phase is described in Section 4.2.1. It includes an initialisation of the domain, the definition of an abstract state, summary, and functions working with them. *Ordering* of abstract states, the *join operator*, and the *widening operator* are defined at the beginning of Chapter 4.

4.2.1 Abstract Domain of the Detection of Atomicity Violations

In this section, at first, it is explained how the abstract domain of this phase is initialised. Then it is described the definition of an abstract state along with functions working with it. In the end, it is described a summary of functions in this phase of the analysis and corresponding functions working with the summary.

Initialisation of the Domain of the Detection of Atomicity Violations

Before analysing each function, i.e., at the beginning of the function `analyse_procedure`, the abstract domain is initialised. The initialisation servers for processing the input file with atomic sequences and storing these sequences into internal data structures in the appropriate format. In the abstract domain, there is a reference for a global data structure `globalData`. The structure contains the following fields:

- **initialised**—it is of a type of a **boolean**. It is used for a determination whether the input file has already been processed.
- **atomicPairs**—it is of a type of a **set of pairs of strings**. It stores pairs of functions that should be called atomically. For instance, $\{(f1\ f2), (f2\ f3)\}$.

The initialisation process is then done the following way. The input file with atomic sequences is read and it is parsed using the *regular expressions*. Each pair of functions that should be called atomically, i.e., any pair of functions in any of the atomic sequences from the input file, is stored into the field `atomicPairs` of the structure `globalData`. Single functions may be also stored into this structure when the atomic sequence contains just one function call. This structure is globally accessible throughout the analysis.

Abstract State of the Domain of the Detection of Atomicity Violations

The abstract state is of a type of `TSet`. `TSet` is a module representing a **set of structures**. This structure has the following fields:

- **firstCall**—it is of a type of a **string**. Captures the first function call within an analysed function. It is used for detection of atomicity violation of pair (a, b) , where a is the last call of a function higher in the *call hierarchy* when calling the analysed function and b is the first function call of the analysed function.
- **lastPair**—is of a type of a **pair of strings**. Captures last two function calls. And it is used for detecting whether this pair violates atomicity. For instance, $(f1\ f2)$.
- **nastedLastCalls**—is of a type of a **list of strings**. Captures the all possible last function calls of the last nested function. It is used for detection of atomicity violation

of pair (a, b) , where a is one of the last calls of the last nested function and b is the analysed function.

- **atomicityViolations**—is of a type of a **set of pairs of strings**. It is used for capturing pairs of function calls that violate atomicity. So it can be reported at the end of the analysis of a function. For instance, $\{(f1\ f2), (f2\ f3)\}$.
- **isInLock**—it is of a type of a **boolean**. Determines whether the current state of a function is inside or outside an atomic block, i.e., it is or it is not under a lock.

The *initial abstract state* is then a set with a single empty element. An empty element is an element where **firstCall** is an empty string, **lastPair** is a pair of two empty strings, **nastedLastCalls** is an empty list, **atomicityViolations** is an empty set, and **isInLock** is **false**.

According to Listings 4.1, there are several functions working with the abstract state. The functions are described below (all of these functions modifies all elements of the abstract state):

- **update_astate_on_function_call**—it is invoked when any function (except a lock or an unlock) is called. When the state is not under a lock, a pair of last two function calls is updated, and it is checked whether this pair (or any pair created from **nastedLastCalls**) violates atomicity. The simplified implementation of the function shows Algorithm 4.4.
- **update_astate_on_lock**—it is invoked when a lock is called. It sets a flag indicating the start of an atomic sequence and it clears the stored last function calls.
- **update_astate_on_unlock**—it is invoked when an unlock is called. It unsets a flag indicating the start of an atomic sequence and it clears the stored last function calls.

Algorithm 4.4: Simplified updating of the abstract state with a called function and checking of atomicity violation

Require: An initialised global data structure *globalData* with a field *atomicPairs* with pairs of functions that should be called atomically

```

1 def update_astate_on_function_call(astate, f):
2   for e  $\in$  astate do
3     if  $\neg(e.isInLock)$  then
4        $(x, y) \leftarrow e.lastPair$ ;
5        $e.lastPair \leftarrow (a, b) \leftarrow (y, f)$ ;
6       if  $(a, b) \in globalData.atomicPairs$  then
7          $e.atomicityViolations \leftarrow Add(e.atomicityViolations, (a, b))$ ;
8       end
9     end
10  end
11  return astate;
12 end

```

Function Summary of the Domain of the Detection of Atomicity Violations

The summary is of a type of **structure**. This structure contains the following fields:

- **firstCalls**—it is of a type of a **list of strings**. It is a list of all possible first function calls of an analysed function.
- **lastCalls**—it is of a type of a **list of strings**. It is a list of all possible last function calls of an analysed function.

Both of the summary fields are used for the purpose of detecting atomicity violating pairs across nested function calls.

According to Listings 4.1 and 4.4, there are some functions working with the summary. The functions are described below:

- **update_astate_on_function_call_with_summary**—it is invoked when the called function has already been analysed. And it performs checks for an atomicity violation of pairs across nested function calls.
- **convert_astate_to_summary**—it is invoked at the end of the analysis of an analysed function. It transforms the abstract state of a given function to the summary. In particular, it derives all the first function calls and all the last function calls within an analysed function from the abstract state.

4.2.2 Reporting of Atomicity Violations

As demonstrates Listing 4.4, at the end of the analysis of a function, atomicity violations within the function are reported. Reporting is achieved by the **Reporting** module implemented in Facebook Infer. For reporting errors using this module, it is necessary to assign an error to the function summary along with a location of the error (a file and line). It is also required to specify a type of the error through the **IssueType** module. The error is then printed to a command line as well as logged to logging files.

Reported atomicity violations are deduced from the abstract state of an analysed function. A simplified reporting process illustrates Algorithm 4.5.

Algorithm 4.5: Reporting of atomicity violations from the abstract state of an analysed function

Input: The abstract state *astate* of an analysed function

```
1 for  $e \in astate$  do
2   for  $(a, b) \in e.atomicityViolations$  do
3      $LogError('„%s”$  and  $„%s”$  should be called atomically.',  $a, b$ );
4   end
5 end
```

4.3 Experimental Verification and Evaluation

This section deals with experimental verification of the analyser and its evaluation. The verification has been performing already during the development of both phases of the analyser (described in Section 3.2). As soon as a part of the analyser has been implemented, it has been verified on appropriate simple programs created for this purpose. In the end, the whole analyser has been successfully verified on smaller programs created for this purpose, as describes Section 4.3.1. Furthermore, Section 4.3.2 shows an experimental evaluation of the analyser on publicly available benchmarks derived from *real-life low-level* programs. Section 4.3.3 then concludes the experimental evaluation and discusses future work.

All the experiments presented in this section are available on attached memory media, see Appendix B. How to run the experiments is described in Appendix C.

4.3.1 Verification on Smaller Hand-Crafted Examples

As it is mentioned above, the verification has been performing already during the development. And both phases of the analyser are verified apart on simple programs written in ANSI C with *PThread* locks.

For **Phase 1** of the analyser, i.e., the *detection of atomic sequences*, it was created several appropriate functions. These functions contain appropriate sequences of function calls inside and outside *atomic blocks*, *not paired lock/unlock calls*, *iteration*, *selection*, and *nested function calls*. The functions were designed in order to check whether single parts of the *abstract domain* work well (i.e., they correspond to the proposal from Chapter 3). In particular, the design aims at verification of the *ordering* operator, the *join* operator, and the *widening* operator. Further, it aims at verification of working with all components of the *abstract state* and the *summary*. It was verified that the result of the analysis of these functions (atomic sequences) is correct, with regard to the proposal from Chapter 3. The analysed functions and the result of **Phase 1** can be seen in Section A.1.

In the case of **Phase 2**, i.e., the *detection of atomicity violations*, the verification process is an analogy of the process for **Phase 1**. But here, there is also verified that reported atomicity violations are reported correctly (based on the atomic sequences from **Phase 1**) with regard to the proposal from Section 3.2.2. The analysed function and the result of **Phase 2** can be seen in Section A.2.

4.3.2 Evaluation on Real-Life Programs

To find out whether the analyser is able to analyse *real programs*, at first, it was tested on a set of slightly more complex student projects from the *Advanced Operating Systems* course. These projects work with threads and synchronisation using *PThreads*.

Later, the analysis was performed on a subset of *real-life low-level concurrent* C programs from a publicly available benchmark. These programs were derived from the Debian GNU Linux distribution. The entire benchmark was originally used for an experimental evaluation of Daniel Kroening’s static deadlock analyser for C/PThreads [14] implemented in the CPROVER framework. For the evaluation, it was used 9 *deadlock-free* programs. The experiments were run on MacBook Pro 2015 with a 2.7 GHz Intel Core i5 processor and

8 GB RAM running the macOS Mojave 10.14.4 operating system. But the running time of the experiments is not relevant because all the experiments were done in less than a few seconds. The results of the experiments are illustrated in Table 4.1. (It is stated the number of lines of code of an analysed program, the number of detected atomic sequences, and the number of detected atomicity violations.)

Program	Lines of Code	Atomic Sequences	Atomicity Violations
alsa-utils 1.1.0	7 735	1	1
c-icap 0.4.2	24 923	11	174
glfw 2.7.9	10 230	9	13
libgroove 4.3.0	7 307	34	294
npth 1.2	1 593	1	26
qrencode 3.4.4	7 006	6	88
rt-tests 0.96	1 795	1	0
signing-party 2.2	1 023	1	1
sslsplit 0.4.11	22 457	18	344

Table 4.1: Experimental results of the analyser on *real-life low-level* programs

As one can see in Table 4.1, in extensive real-life programs, there are reported quite a lot of atomicity violations. Many of them are probably *false alarms*. Determination, whether the reported atomicity violations in such real-life programs are actual errors may be challenging and it goes beyond the scope of this thesis. However, a result of the analyser can be used as an input for *dynamic analysis* which can be able to ascertain if reported atomicity violations are real errors.

4.3.3 Summary of the Evaluation and Future Work

The correctness of the analyser was successfully verified on smaller *hand-crafted* example programs. However, when analysing extensive *real-life* programs, many *false alarms* are usually reported. To reduce the number of false alarms, it seems promising to work with *nested locks*, consider the *different locks* used, consider extensions for *contracts for concurrency* introduced in Section 2.3, i.e., consider *function parameters* and *contextual information* of function calls. Another option for the reduction of false alarms is to working with sets instead of pairs when looking for functions that should be called atomically. The considerable number of false alarms are caused by atomic sequences that contain only one function call and by atomic sequences which do not have to be called atomically always. The solution could be ignoring *atomic blocks* where there is just one function call, and consider atomic sequences only if they appear in an atomic block more than, e.g., three times.

Another issue is that analysing of more complex programs with extensive *control structures* and a lot of function calls inside atomic blocks is very time and memory consuming and the analysis either fails or it runs an extremely long time. The solution of this issue could be replacement working with the **A** and **B** sequences in **Phase 1** of the analysis by working with sets. This replacement would be at the expense of precision but it would definitely gain speed.

Chapter 5

Conclusion

This thesis explains the principles of *static analysis* and *abstract interpretation*. Further, it describes a static analysis framework that uses abstract interpretation — *Facebook Infer* — its features, architecture, and existing analysers implemented in this tool. Furthermore, it introduces and defines the concept of *contracts for concurrency*. The major part of the thesis aims at the proposal of a static analyser for detecting *atomicity violations* — *Atomer* — and its implementation as an extension for Facebook Infer. Lastly, it is described the experimental evaluation of the implemented analyser and it is discussed the future work.

The analyser works on the level of *sequences of function calls*. The proposed solution is based on the assumption that sequences executed *once atomically* should probably be executed *always atomically*. It is also inspired by the concept of contracts for concurrency. *Atomer* is divided into two phases of the analysis: **Phase 1** – detection of *atomic sequences*, i.e., detection of function calls executed atomically; and **Phase 2** – detection of *atomicity violations*, i.e., detection of violations of the atomic sequences obtained from **Phase 1**. The implementation targets C/C++ programs that use *PThread* locks.

Atomer has been successfully verified on smaller *hand-crafted* programs. Moreover, it has been experimentally evaluated on publicly available benchmarks derived from *real-life low-level* programs from the Debian distribution. The analysis has been performed on 9 low-level programs with a couple of thousands of lines of code. It has been found out that *Atomer* is able to analyse such extensive real-life programs, but in such cases, it reports quite a lot of *false alarms*. Anyway, a result of the analyser may be used as an input for *dynamic analysis* which can be able to determine whether reported atomicity violations are real errors.

Atomer shows the potential for further improvements. The future work will focus mainly on increasing the accuracy of used methods using, e.g., considering *nested locks*, *different locks* used, *function parameters*, etc. The future work will also focus on enhancing the *scalability* because *Atomer* fails to analyse more extensive and complex programs. Further, it would be interesting to extend the analysis for *other types of locks* for synchronisation of concurrent threads/processes and testing the analysis on other real-world programs.

The code of *Atomer* is available on GitHub as an *open-source repository*. It would be nice to achieve a merge of the implementation of *Atomer* into the **master** branch of Facebook Infer’s repository. The preliminary results of the thesis were published and presented in paper [13] for the student conference *Excel@FIT*, where it won an award in two categories.

Bibliography

- [1] Allen, F. E.: Control Flow Analysis. In *Proceedings of a Symposium on Compiler Optimization*. Urbana-Champaign, Illinois: ACM, New York, NY, USA. 1970. pp. 1 – 19. doi:10.1145/800028.808479.
- [2] Blackshear, S.; Distefano, D.; Villard, J.: Building your own compositional static analyzer with Infer.AI [online]. PLDI'17. 2017 [cit. 2019-05-03]. Retrieved from: <https://fbinfer.com/downloads/pldi17-infer-ai-tutorial.pdf>
- [3] Blackshear, S.; Gorogiannis, N.; O'Hearn, P. W.; Sergey, I.: RacerD: Compositional Static Race Detection. *Proceedings of ACM Programming Languages*. vol. 2, no. OOPSLA'18. October 2018: pp. 144:1 – 144:28. ISSN 2475-1421. doi:10.1145/3276514.
- [4] Blackshear, S.; O'Hearn, P. W.: Open-sourcing RacerD: Fast static race detection at scale [online]. 2017-10-19 [cit. 2019-05-03]. Retrieved from: <https://code.fb.com/android/open-sourcing-racerd-fast-static-race-detection-at-scale>
- [5] Calcagno, C.; Distefano, D.; O'Hearn, P. W.; Yang, H.: Compositional Shape Analysis by Means of Bi-abduction. In *Proceedings of the 36th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL'09. Savannah, GA, USA: ACM, New York, NY, USA. January 2009. ISBN 978-1-60558-379-2. pp. 289 – 300. doi:10.1145/1480881.1480917.
- [6] Cousot, P.: Abstract Interpretation in a Nutshell [online]. [cit. 2019-05-02]. Retrieved from: <https://www.di.ens.fr/~cousot/AI/IntroAbsInt.html>
- [7] Cousot, P.: Abstract Interpretation [online]. 2008-08-05 [cit. 2019-05-02]. Retrieved from: <https://www.di.ens.fr/~cousot/AI>
- [8] Cousot, P.: Abstract Interpretation Based Formal Methods and Future Challenges, invited paper. In « *Informatics — 10 Years Back, 10 Years Ahead* », *Lecture Notes in Computer Science*, vol. 2000, edited by R. Wilhelm. Springer-Verlag. March 2001. pp. 138 – 156. doi:10.1007/3-540-44577-3_10.
- [9] Cousot, P.; Cousot, R.: Abstract interpretation: a unified lattice model for static analysis of programs by construction or approximation of fixpoints. In *Conference Record of the Fourth Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL'77. Los Angeles, California: ACM Press, New York, NY. 1977. pp. 238 – 252. doi:10.1145/512950.512973.

- [10] Cousot, P.; Cousot, R.: Comparing the Galois Connection and Widening/Narrowing Approaches to Abstract Interpretation, invited paper. In *Proceedings of the International Workshop Programming Language Implementation and Logic Programming*. PLILP'92, edited by M. Bruynooghe; M. Wirsing. Leuven, Belgium, 13–17 August 1992, Lecture Notes in Computer Science 631. Springer-Verlag, Berlin, Germany. January 1992. pp. 269–295. doi:10.1007/3-540-55844-6_101.
- [11] Dias, R. J.; Ferreira, C.; Fiedor, J.; Lourenço, J. M.; Smrčka, A.; Sousa, D. G.; Vojnar, T.: Verifying Concurrent Programs Using Contracts. In *2017 IEEE International Conference on Software Testing, Verification and Validation*. ICST'17. Tokyo, Japan: IEEE. March 2017. ISBN 9781509060313. pp. 196–206. doi:10.1109/ICST.2017.25.
- [12] Gorogiannis, N.; O'Hearn, P. W.; Sergey, I.: A True Positives Theorem for a Static Race Detector. *Proceedings of ACM Programming Languages*. vol. 3, no. POPL'19. January 2019: pp. 57:1–57:29. ISSN 2475-1421. doi:10.1145/3290370.
- [13] Harmim, D.; Marin, V.; Pavela, O.: Scalable Static Analysis Using Facebook Infer. In *Excel@FIT*. Brno University of Technology, Faculty of Information Technology. 2019.
- [14] Kroening, D.; Poetzl, D.; Schrammel, P.; Wachter, B.: Sound static deadlock analysis for C/Pthreads. In *Proceedings of the 31st IEEE/ACM International Conference on Automated Software Engineering*. ASE'16. Singapore, Singapore: ACM, New York, NY, USA. September 2016. ISBN 978-1-4503-3845-5. pp. 379–390. doi:10.1145/2970276.2970309.
- [15] Lengál, O.; Vojnar, T.: Abstract Interpretation. Lecture Notes in Formal Analysis and Verification. Brno University of Technology, Faculty of Information Technology. 2018.
- [16] Marcin, V.: Static Analysis of Concurrency Problems in the Facebook Infer Tool. Project practice. Brno University of Technology, Faculty of Information Technology. 2018.
- [17] Meyer, B.: Applying „Design by Contract“. *Computer*. vol. 25, no. 10. October 1992: pp. 40–51. ISSN 0018-9162. doi:10.1109/2.161279.
- [18] Minsky, Y.; Madhavapeddy, A.; Hickey, J.: *Real world OCaml*. Sebastopol, CA: O'Reilly Media. first edition. 2013. ISBN 144932391X.
- [19] Møller, A.; Schwartzbach, I. M.: *Static Program Analysis*. Department of Computer Science, Aarhus University. October 2018.
- [20] Nielson, F.; Nielson, R. H.; Hankin, C.: *Principles of Program Analysis*. Berlin: Springer-Verlag. 2005. ISBN 3-540-65410-0.
- [21] Reps, T.; Horwitz, S.; Sagiv, M.: Precise Interprocedural Dataflow Analysis via Graph Reachability. In *Proceedings of the 22Nd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. POPL'95. ACM, New York, NY, USA. 1995. ISBN 0-89791-692-1. pp. 49–61. doi:10.1145/199448.199462.

- [22] Sharir, M.; Pnueli, A.: Two approaches to interprocedural data flow analysis. In *Program Flow Analysis: Theory and Applications*, edited by S. S. Muchnick; N. D. Jones. chapter 7. Prentice Hall Professional Technical Reference. 1981. ISBN 0137296819. pp. 189–211.
- [23] Sousa, D. G.; Dias, R. J.; Ferreira, C.; Lourenço, J. M.: Preventing Atomicity Violations with Contracts. *CoRR*. vol. abs/1505.02951. 2015. [1505.02951](#).
- [24] Vojnar, T.: Different Approaches to Formal Verification and Analysis. Lecture Notes in Formal Analysis and Verification. Brno University of Technology, Faculty of Information Technology. 2018.
- [25] Vojnar, T.: Lattices and Fixpoints for Symbolic Model Checking. Lecture Notes in Formal Analysis and Verification. Brno University of Technology, Faculty of Information Technology. 2018.
- [26] Yi, K.: Inferbo: Infer-based buffer overrun analyzer [online]. 2017-02-06 [cit. 2019-05-04]. Retrieved from: <https://research.fb.com/inferbo-infer-based-buffer-overrun-analyzer>

List of Appendices

A	Experimental Verification Results	39
A.1	Detection of Atomic Sequences	39
A.2	Detection of Atomicity Violations	42
B	Contents of Attached Memory Media	45
C	Installation and User Manual	46
C.1	Installation Manual	46
C.2	User Manual	46

Appendix A

Experimental Verification Results

This appendix illustrates results of experimental verification of the implemented analyser. The verification process and the results, in general, are in more detail discussed in Section 4.3. Section A.1 shows experimental results of **Phase 1** of the analyser, i.e., the *detection of atomic sequences*. Section A.2 then shows experimental results of **Phase 2**, i.e., the *detection of atomicity violations*.

Both sections demonstrate the analysis on programs written in ANSI C and assume *PThread* locks and the existence of an initialised global variable `lock` of a type `pthread_mutex_t`.

Moreover, the below experiments are available on attached memory media, see Appendix B. Instructions how to run these experiments can be found in Appendix C.

A.1 Detection of Atomic Sequences

For the verification of the detection of atomic sequences are used functions defined in Listing A.1. The result of this detection, i.e., *sequences of function that should be called atomically*, is shown in Listing A.2.

```
1 void f1(void) {}
2 void f2(void) {}
3 void f3(void) {}
4 void f4(void) {}
5 void f5(void) {}
6 void ff(void) { f1(); f2(); }
7
8 void test1(void)
9 {
10     f1(); f1();
11
12     pthread_mutex_lock(&lock);
13     f1(); f1(); f2();
14     pthread_mutex_unlock(&lock);
15
16     f1(); f1();
```

```

17
18     pthread_mutex_lock(&lock);
19     f1(); f3();
20     pthread_mutex_unlock(&lock);
21
22     f1();
23
24     pthread_mutex_lock(&lock);
25     f1(); f3(); f3();
26     pthread_mutex_unlock(&lock);
27 }
28
29 void test2(void)
30 {
31     f1(); f1();
32
33     pthread_mutex_lock(&lock);
34     f1(); f1(); f2();
35     pthread_mutex_unlock(&lock);
36
37     f3(); f3();
38
39     pthread_mutex_lock(&lock);
40     f1(); f4(); f4();
41     pthread_mutex_unlock(&lock);
42 }
43
44 void test_only_lock(void)
45 {
46     pthread_mutex_lock(&lock);
47     f1();
48 }
49
50 void test_only_unlock(void)
51 {
52     f2();
53     pthread_mutex_unlock(&lock);
54 }
55
56 void test_iteration(void)
57 {
58     int c;
59     f1(); f2();
60     while (c > 0) { f3(); f5(); }
61
62     pthread_mutex_lock(&lock);
63     f1(); f2();
64     pthread_mutex_unlock(&lock);

```

```

65
66     pthread_mutex_lock(&lock);
67     while (c > 0) f3();
68     pthread_mutex_unlock(&lock);
69
70     f4();
71 }
72
73 void test_selection(void)
74 {
75     int c;
76     f1(); f2();
77     if (c > 0) { f3(); f5(); }
78     else
79     {
80         pthread_mutex_lock(&lock);
81         f1();
82         pthread_mutex_unlock(&lock);
83     }
84
85     pthread_mutex_lock(&lock);
86     f2();
87     if (c > 42) f3();
88     else if (c > 0) f4();
89     pthread_mutex_unlock(&lock);
90
91     f4();
92 }
93
94 void test_nested(void)
95 {
96     pthread_mutex_lock(&lock);
97     ff(); f3();
98     pthread_mutex_unlock(&lock);
99
100    pthread_mutex_lock(&lock);
101    f4(); f5(); ff();
102    pthread_mutex_unlock(&lock);
103 }

```

Listing A.1: Functions to be analysed for the detection of atomic sequences

```

f1:␣
f2:␣
f3:␣
f4:␣
f5:␣
ff:␣
test1:␣(f1␣f2)␣(f1␣f3)

```

```

test2:_(f1_f2)_(f1_f4)
test_only_lock:_(f1)
test_only_unlock:_
test_iteration:_(f1_f2)_(f3)
test_selection:_(f1)_(f2)_(f2_f3)_(f2_f4)
test_nested:_(ff_f1_f2_f3)_(f4_f5_ff_f1_f2)

```

Listing A.2: The result of **Phase 1** (atomic sequences) of the analysis of functions from Listing A.1

A.2 Detection of Atomicity Violations

For the verification of the detection of atomicity violations are used functions defined in Listing A.3. The result of the detection of atomic sequences is shown in Listing A.4. The result of the detection of atomicity violations, i.e., *functions that should be called atomically but they are not*, is stated in Listing A.3 using **comments**.

```

1 void f1(void) {}
2 void f2(void) {}
3 void f3(void) {}
4 void f4(void) {}
5 void g(void) {}
6 void ff(void) { f3(); f1(); f4(); } // (f3 f1) (f1 f4)
7
8 void atomic_sequences(void)
9 {
10     pthread_mutex_lock(&lock);
11     f1(); f2(); f3();
12     pthread_mutex_unlock(&lock);
13
14     pthread_mutex_lock(&lock);
15     f4(); f2();
16     pthread_mutex_unlock(&lock);
17
18     pthread_mutex_lock(&lock);
19     f1(); f3();
20     pthread_mutex_unlock(&lock);
21
22     pthread_mutex_lock(&lock);
23     ff(); f3();
24     pthread_mutex_unlock(&lock);
25 }
26
27 void test1(void)
28 {
29     f1(); f2(); g(); // (f1 f2)
30     f1(); g(); f2(); g();

```



```

31     f1(); f1(); f2(); g(); // (f1 f2)
32     f1(); f2(); f3(); g(); // (f1 f2) (f2 f3)
33     f1(); g(); f2(); g(); f3();
34 }
35
36 void test2(void)
37 {
38     f4(); f2(); g(); // (f4 f2)
39     f2(); f4(); g();
40
41     f4();
42     pthread_mutex_lock(&lock);
43     pthread_mutex_unlock(&lock);
44     f2();
45     g();
46
47     f3(); f4();
48 }
49
50 void test_only_lock(void)
51 {
52     pthread_mutex_lock(&lock);
53     f1(); f2();
54 }
55
56 void test_only_unlock(void)
57 {
58     f1(); f2(); // (f1 f2)
59     pthread_mutex_unlock(&lock);
60 }
61
62 void test_iteration(void)
63 {
64     int c;
65     while (c > 0) { f1(); f2(); } // (f1 f2)
66
67     f1();
68     while (c > 0) f2(); // (f1 f2)
69     f3(); // (f1 f3) (f2 f3)
70
71     for (; c > 0; f1()) f3(); // (f3 f1) (f1 f3)
72
73     pthread_mutex_lock(&lock);
74     f1(); f2();
75     while (c > 0) { f2(); f3(); }
76     pthread_mutex_unlock(&lock);
77 }
78

```

```

79 void test_selection(void)
80 {
81     int c;
82     f1();
83     if (c > 0) f2(); // (f1 f2)
84     else f3(); // (f1 f3)
85     f3(); // (f2 f3)
86
87     g();
88
89     f4();
90     if (c > 42) f4();
91     else if (c > 0) f2(); // (f4 f2)
92     f2(); // (f4 f2)
93 }
94
95 void test_nested(void)
96 {
97     pthread_mutex_lock(&lock);
98     ff();
99     pthread_mutex_unlock(&lock);
100
101     ff(); g(); // (ff f3)
102     ff(); f2(); // (ff f3) (f4 f2)
103 }

```

Listing A.3: Functions to be analysed for the detection of atomic sequences and for the subsequent detection of atomicity violations (detected atomicity violations are stated in **comments**)

```

f1:␣
f2:␣
f3:␣
f4:␣
g:␣
ff:␣
atomic_sequences:␣(f1␣f2␣f3)␣(f4␣f2)␣(f1␣f3)␣(ff␣f3␣f1␣f4)
test1:␣
test2:␣
test_only_lock:␣(f1␣f2)
test_only_unlock:␣
test_iteration:␣(f1␣f2)␣(f1␣f2␣f3)
test_selection:␣
test_nested:␣(ff␣f3␣f1␣f4)

```

Listing A.4: The result of **Phase 1** (atomic sequences) of the analysis of functions from Listing A.3

Appendix B

Contents of Attached Memory Media

[[**TODO**]]

Appendix C

Installation and User Manual

[[TODO]]

C.1 Installation Manual

[[TODO]]

C.2 User Manual

[[TODO]]